



# Module 2: Exploratory Data Analysis

**Part 2:** Two Dimensional Data



# Exercise 1

For this exercise, you are given a Series of random numbers creatively names **random\_numbers**. For the first exercise please do the following:

- 1 Remove all the numbers less than 10
- 2 Sort the series
- 3 Calculate the Tukey 5 number summary for this dataset
- 4 Count the number of even and odd numbers
- 5 Find the five largest and 5 smallest numbers in the series



# Exercise 1

```
#Filter the Series
```

```
random_numbers = random_numbers[random_numbers >= 10]
```

```
#Sort the Series
```

```
random_numbers.sort_values(inplace=True)
```

```
#Calculate the Tukey 5 Number Summary
```

```
random_numbers.describe()
```

```
#Count the number of even and odd numbers
```

```
even_numbers = random_numbers[random_numbers % 2 == 0].count()
```

```
odd_numbers = random_numbers[random_numbers % 2 != 0].count()
```

```
print( "Even numbers: " + str(even_numbers))
```

```
print( "Odd numbers: " + str(odd_numbers))
```

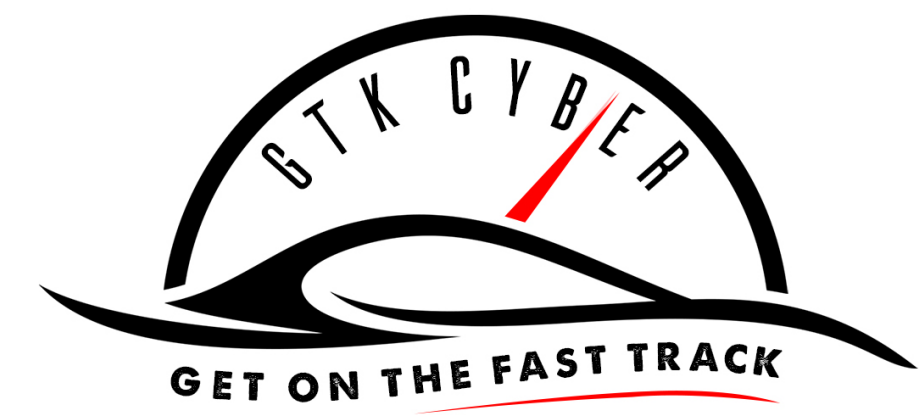
```
#Find the five largest and smallest numbers
```

```
print( "Smallest Numbers:")
```

```
print( random_numbers.head(5) )
```

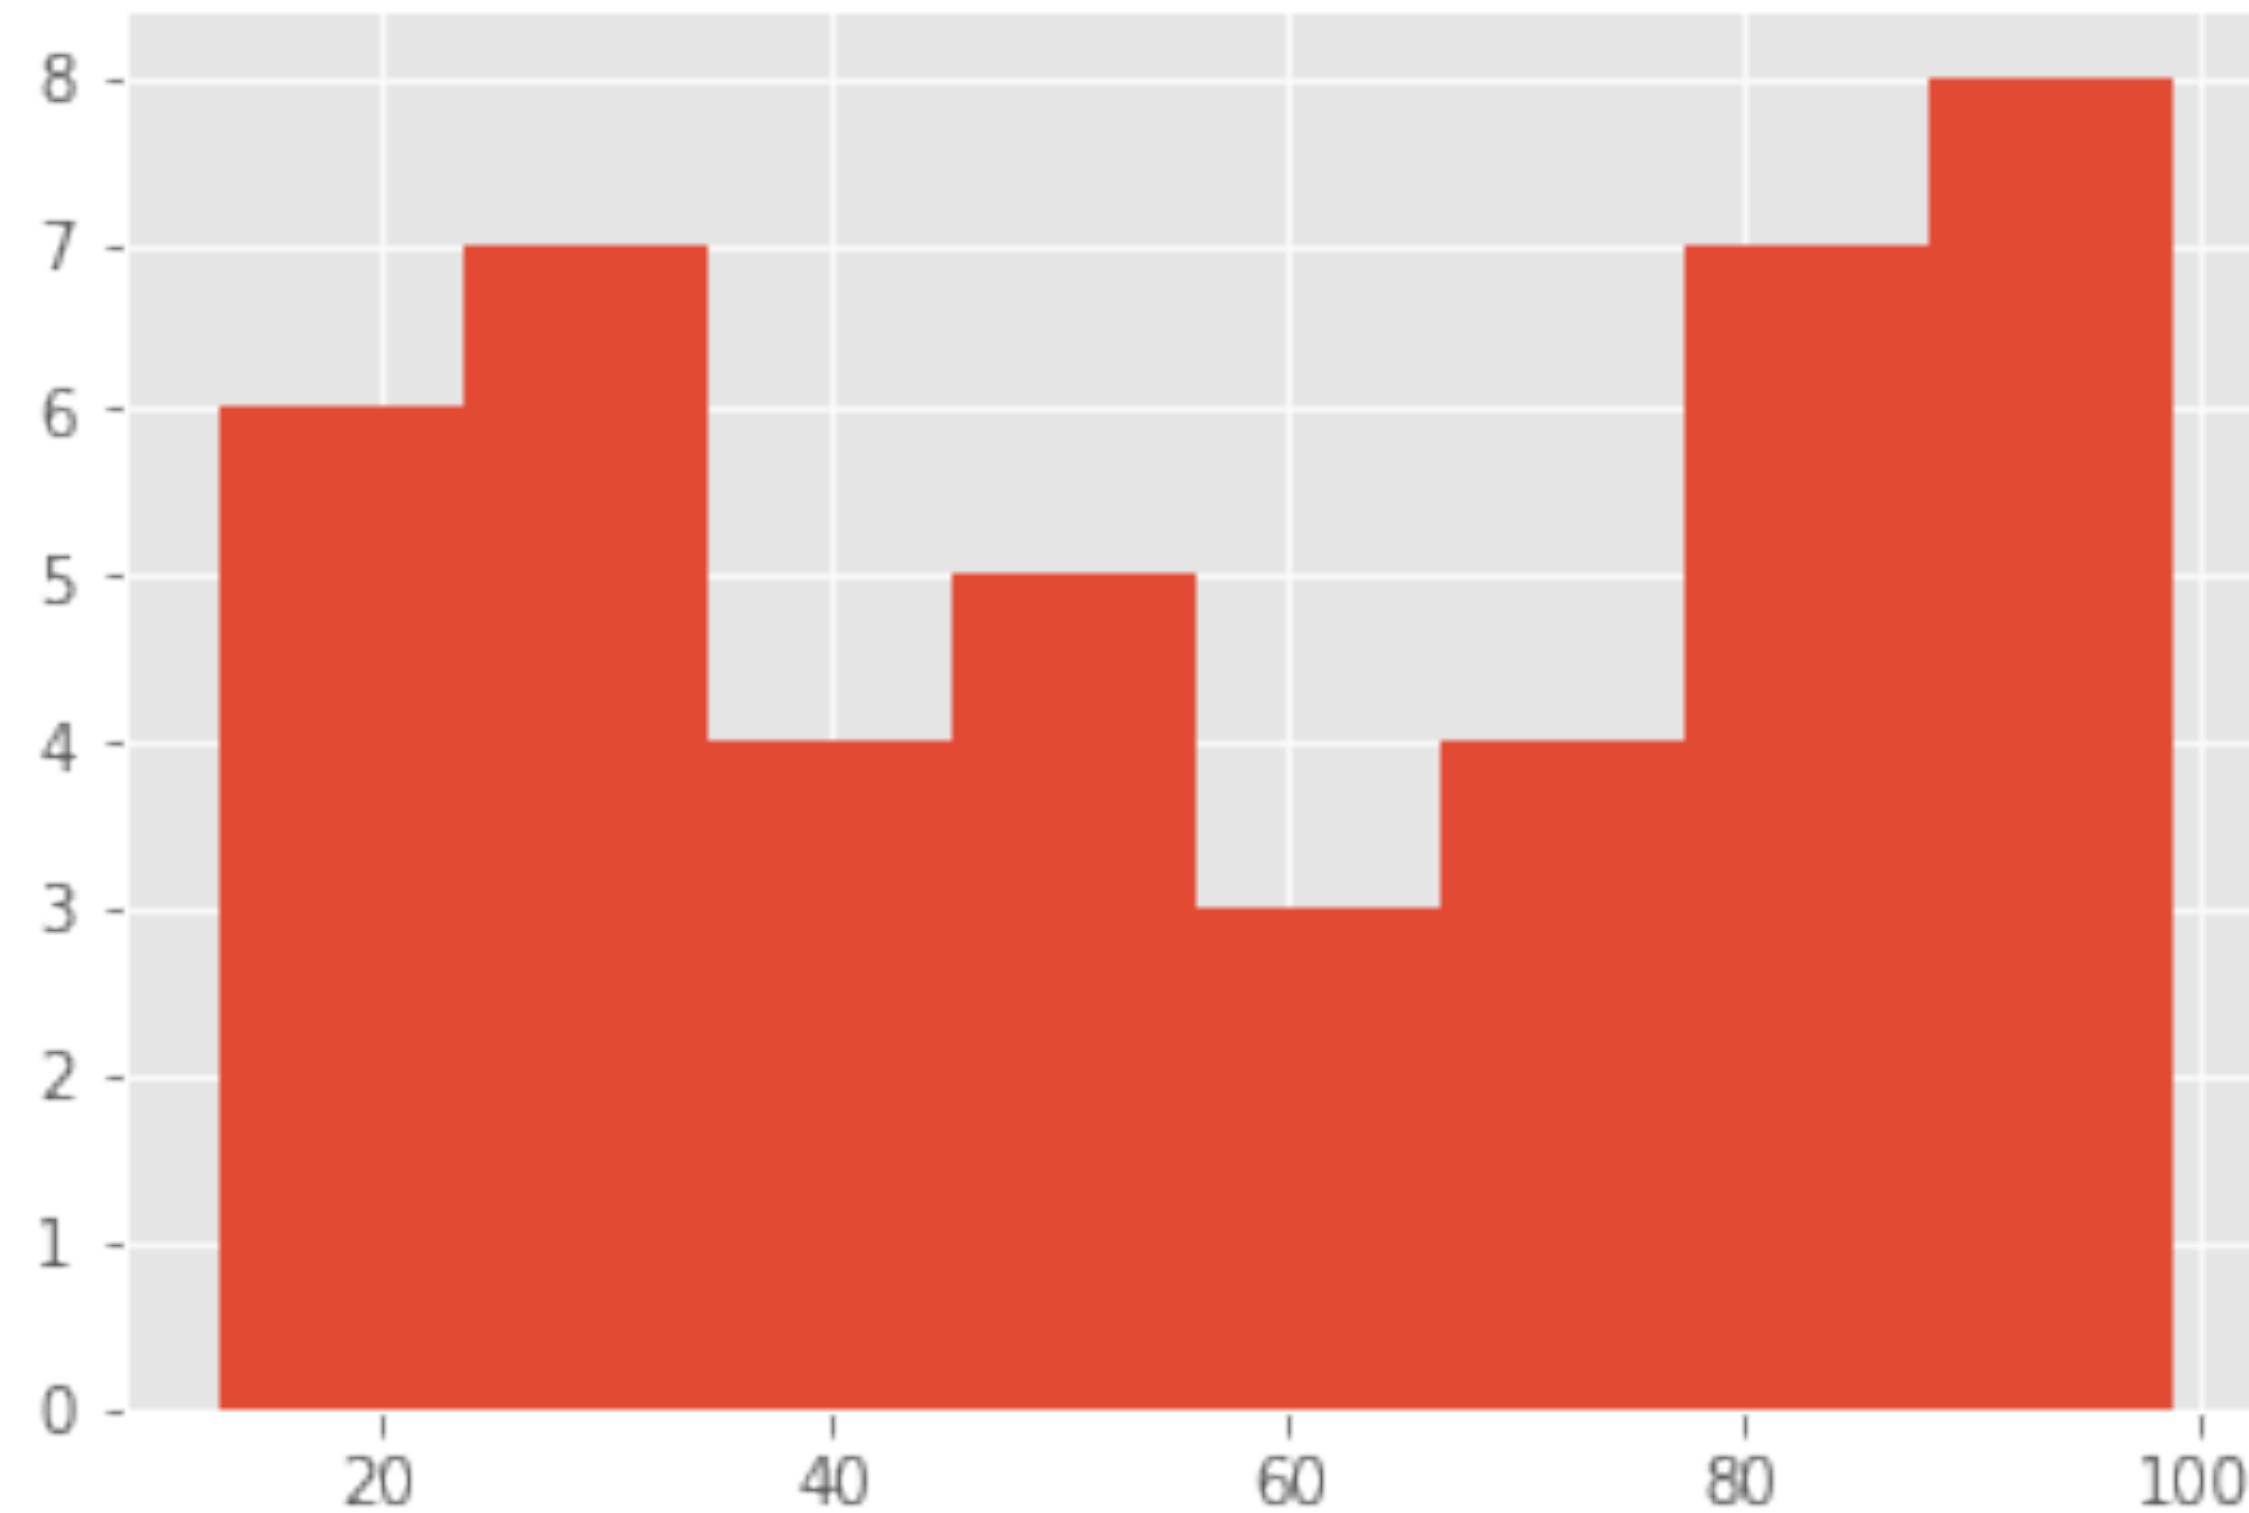
```
print( "Largest Numbers:")
```

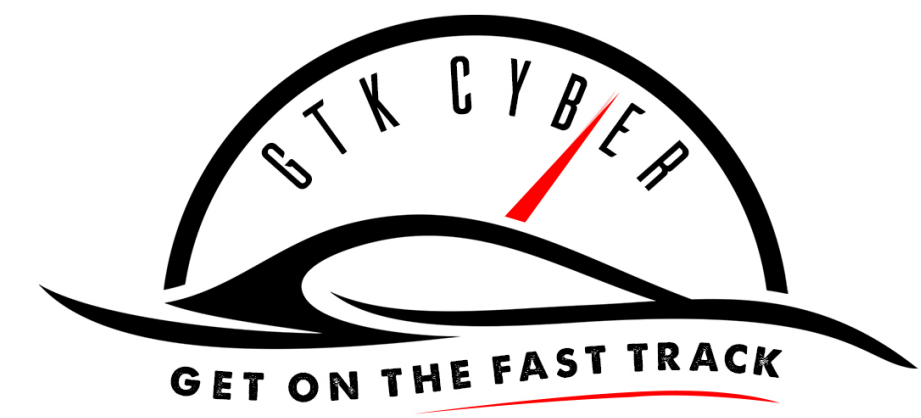
```
print( random_numbers.tail(5) )
```



# Exercise 2

```
random_numbers.hist(bins=8)
```



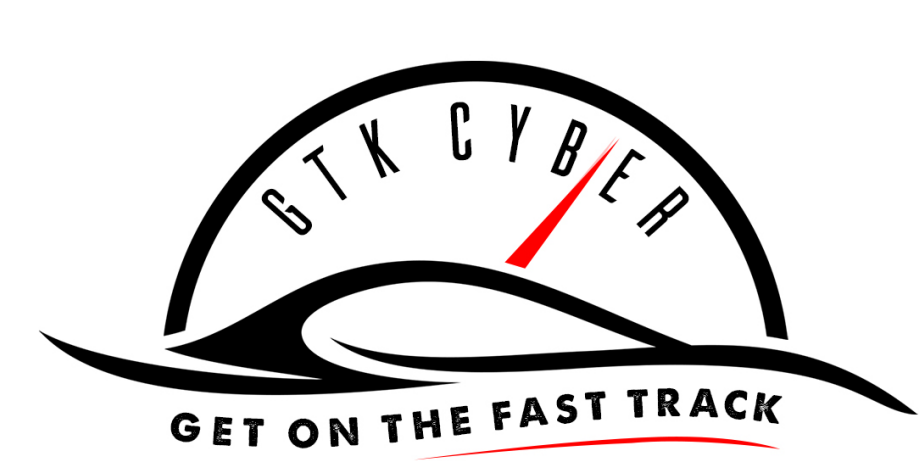


# Exercise 3

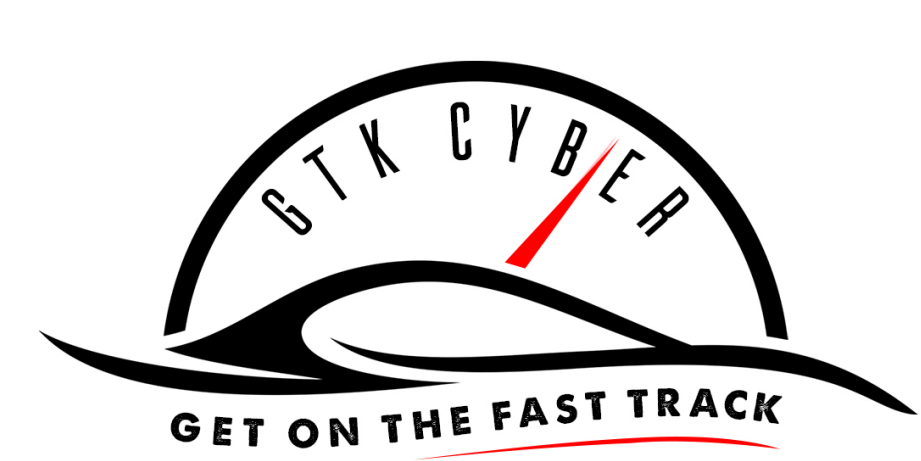
```
phone_number_series = pd.Series(phone_numbers)

area_codes = phone_number_series.str.slice(1,4)
area_codes2 = phone_number_series.str.extract( '\((\d{3})\)')

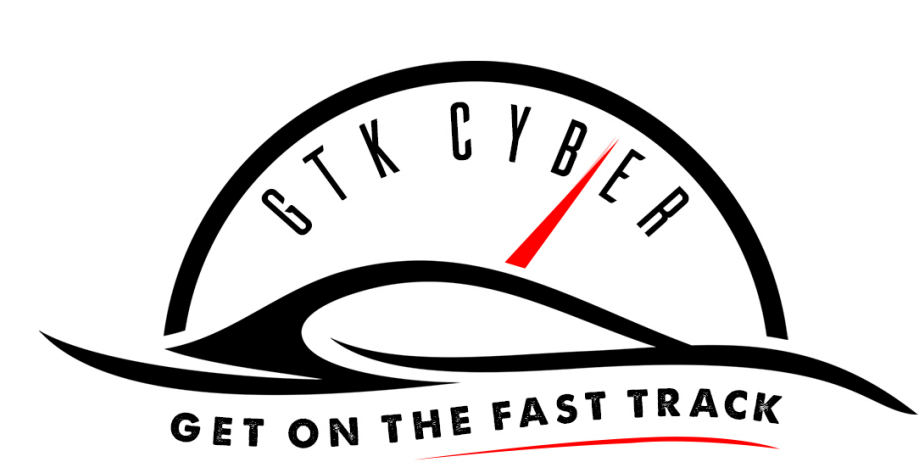
area_codes.value_counts()
```



# The Data Frame



```
df = pd.DataFrame( <data>, <index>, <column_names> )
```



# CSV

```
df = pd.read_csv( <file> )
```

```
df = pd.read_csv( <url> )
```



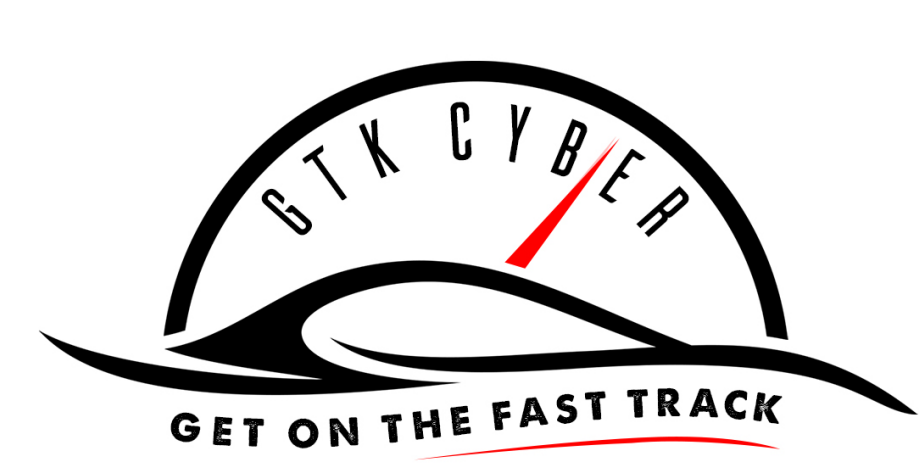


# Excel

```
df = pd.read_excel( <file>, sheetname=<sheetname> )
```



```
[  
  {  
    "changed": "2015-04-03 00:00:00",  
    "created": "2014-04-10 00:00:00",  
    "dnssec": "False",  
    "expires": "2016-04-10 00:00:00",  
    "isMalicious": 0,  
    "url": "nuteczki.com"  
  },  
  ...  
]
```



# JSON

```
df = pd.read_json( <file> )
```

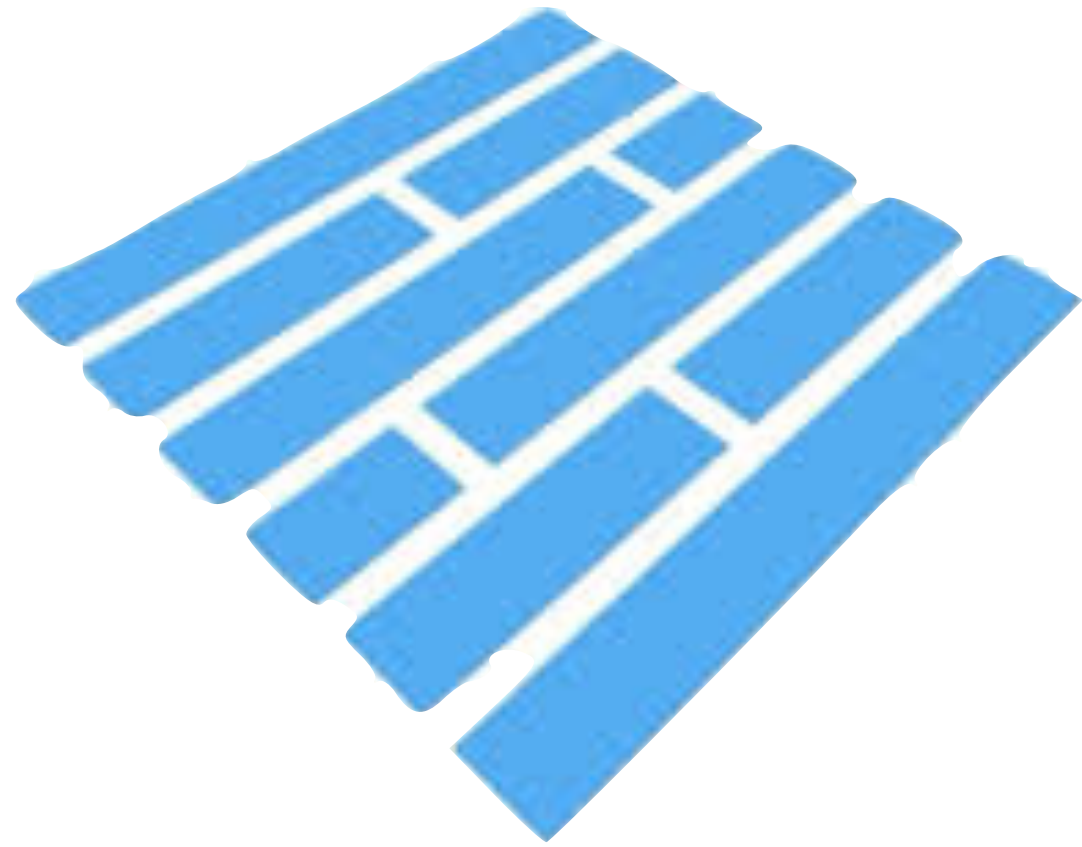
or

```
df = pd.read_json( <url> )
```



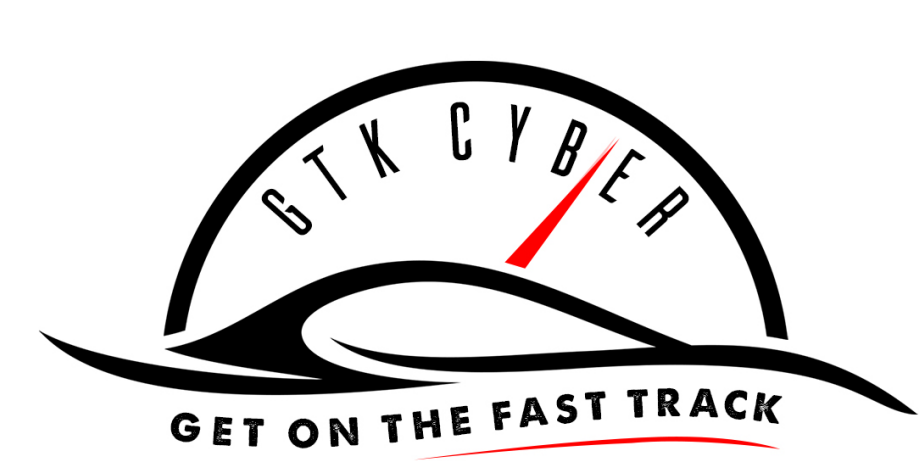
# From a Database

```
df = pd.read_sql( <query>, <connection> )
```



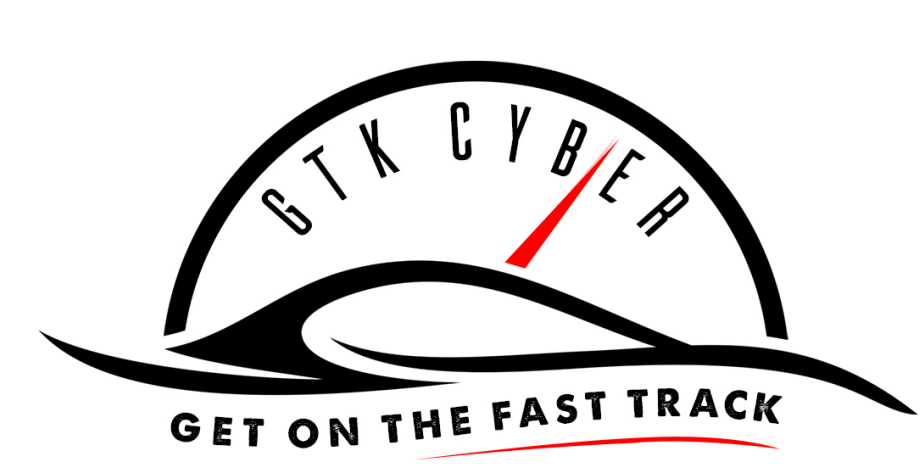
# Parquet

```
df = pd.read_parquet( <file> )
```



# HTML

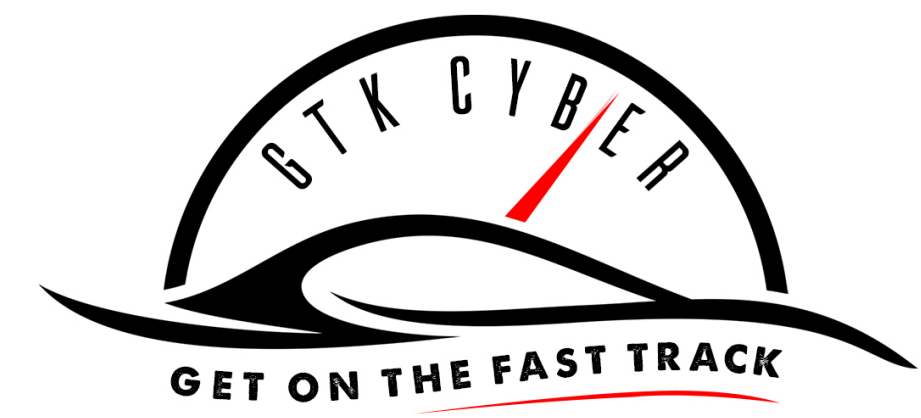
```
df = pd.read_html( <source> )
```



# XML

```
import requests
```

```
user_agent_url = 'http://www.user-agents.org/allagents.xml'  
xml_data = requests.get(user_agent_url).content
```



# XML

```
import xml.etree.ElementTree as ET

class XML2DataFrame:

    def __init__(self, xml_data):
        self.root = ET.XML(xml_data)

    def parse_root(self, root):
        return [self.parse_element(child) for child in iter(root)]

    def parse_element(self, element, parsed=None):
        if parsed is None:
            parsed = dict()
        for key in element.keys():
            parsed[key] = element.attrib.get(key)
        if element.text:
            parsed[element.tag] = element.text
        for child in list(element):
            self.parse_element(child, parsed)
        return parsed

    def process_data(self):
        structure_data = self.parse_root(self.root)
        return pd.DataFrame(structure_data)

xml2df = XML2DataFrame(xml_data)
xml_dataframe = xml2df.process_data()
```

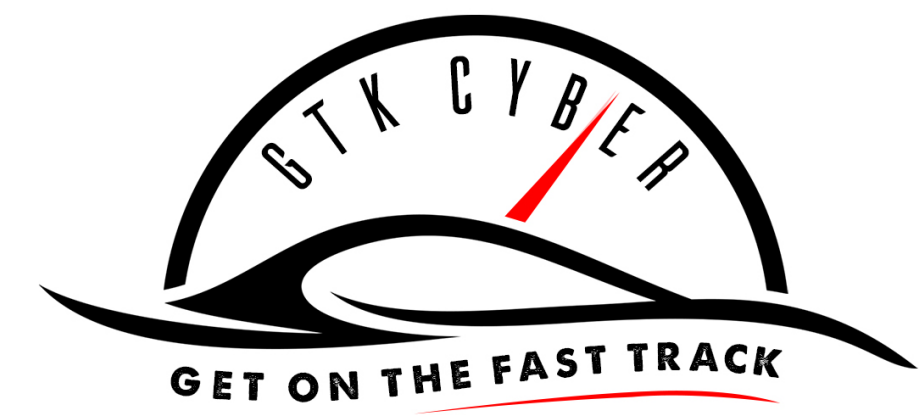




# Log Files...



070823	21:00:32	1	Connect	root@localhost on test1
070823	21:00:48	1	Query	show tables
070823	21:00:56	1	Query	select * from category
070917	16:29:01	21	Query	select * from location
070917	16:29:12	21	Query	select * from location where id = 1 LIMIT 1



```
070823 21:00:32      1 Connect      root@localhost on test1
070823 21:00:48      1 Query        show tables
070823 21:00:56      1 Query        select * from category
070917 16:29:01     21 Query        select * from location
070917 16:29:12     21 Query        select * from location where id = 1 LIMIT 1
```

```
logdf = pd.read_table('../data/mysql.log', names=[ 'raw' ])
```

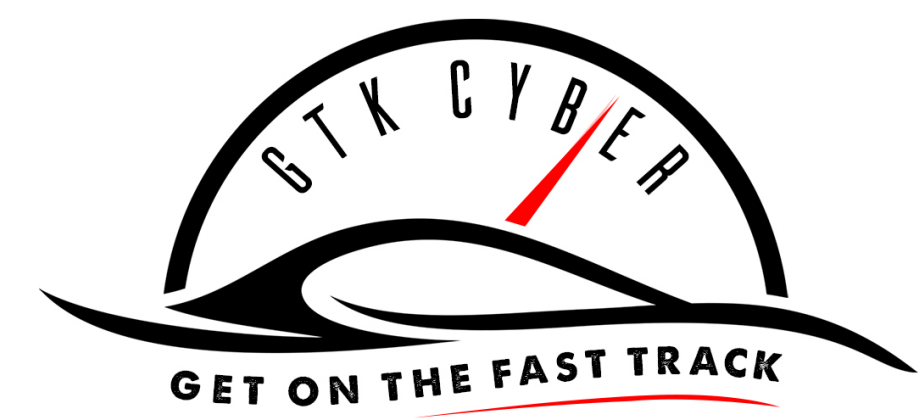
	raw
0	070823 21:00:32 1 Connect root@local...
1	070823 21:00:48 1 Query show tables
2	070823 21:00:56 1 Query select * f...
3	070917 16:29:01 21 Query select * f...
4	070917 16:29:12 21 Query select * f...



```
logdf = pd.read_table('../data/mysql.log', names=[ 'raw' ])
logdf[ 'raw' ].str.extract( ' (\\d{6}\\s\\d{2}:\\d{2}:\\d{2})\\s+(\\d+)\\s(\\S+)\\s(\\S+)' , expand=False)
```

	0	1	2	3
0	070823 21:00:32	1	Connect	root@localhost on test1
1	070823 21:00:48	1	Query	show tables
2	070823 21:00:56	1	Query	select * from category
3	070917 16:29:01	21	Query	select * from location
4	070917 16:29:12	21	Query	select * from location where id = 1 LIMIT 1





```
logdf = pd.read_table('../data/mysql.log', names=[ 'raw' ])
logdf[ 'raw' ].str.extract( '(?P<date>\d{6}\s\d{2}:\d{2}:\d{2})\s+(?P<PID>\d+)\s(?P<Action>S+)\s(?P<Query>.+)', expand=False)
```

	Date	PID	Action	Query
0	070823 21:00:32	1	Connect	root@localhost on test1
1	070823 21:00:48	1	Query	show tables
2	070823 21:00:56	1	Query	select * from category
3	070917 16:29:01	21	Query	select * from location
4	070917 16:29:12	21	Query	select * from location where id = 1 LIMIT 1



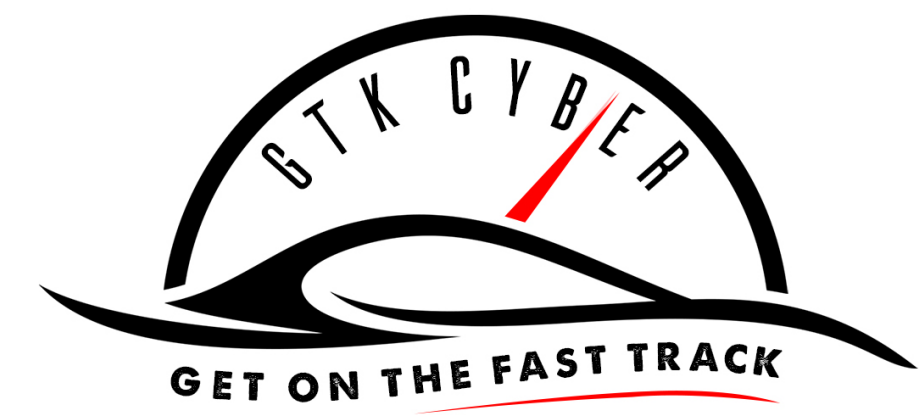
# Web Server Logs





# Web Server Logs

```
195.154.46.135 - - [25/Oct/2015:04:11:25 +0100] "GET /linux/  
doing-pxe-without-dhcp-control HTTP/1.1" 200 24323 "http://  
howto.basjes.nl/" "Mozilla/5.0 (Windows NT 5.1; rv:35.0)  
Gecko/20100101 Firefox/35.0"
```

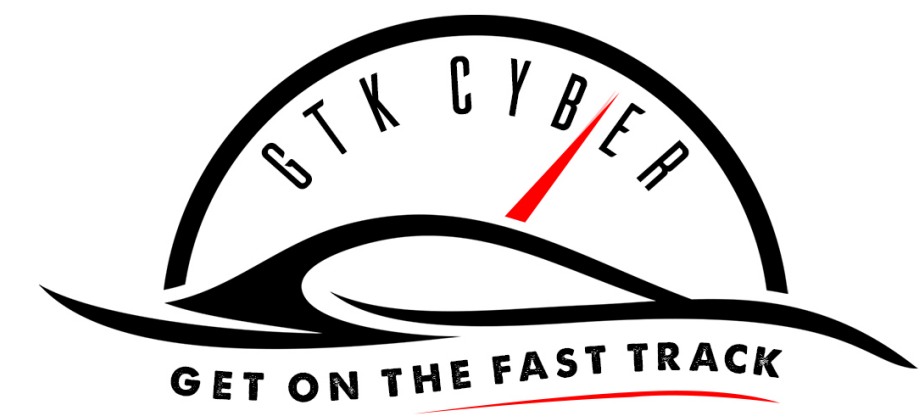


# Web Server Logs

```
195.154.46.135 - - [25/Oct/2015:04:11:25 +0100] "GET /linux/  
doing-pxe-without-dhcp-control HTTP/1.1" 200 24323 "http://  
howto.basjes.nl/" "Mozilla/5.0 (Windows NT 5.1; rv:35.0)  
Gecko/20100101 Firefox/35.0"
```

```
import apache_log_parser  
line_parser = apache_log_parser.make_parser("%h %l %u %t  
\"%r\" %>s %b \"%{Referer}i\" \"%{User-agent}i\"")
```



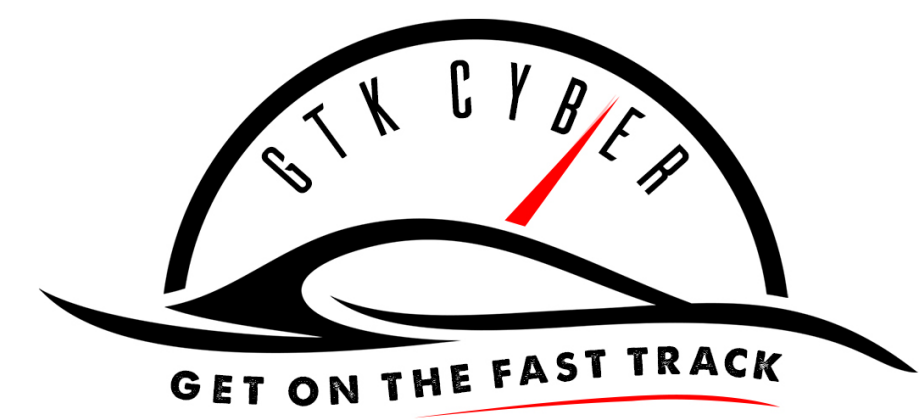


# Web Server Logs

```
import apache_log_parser
line_parser = apache_log_parser.make_parser("%h %l %u %t
\"%r\" %>s %b \"%{Referer}i\" \"%{User-agent}i\"")

server_log = open("../data/hackers-access.httpd", "r")
parsed_server_data = []
for line in server_log:
    data = {}
    data = line_parser(line)
    parsed_server_data.append( data )

server_df = pd.DataFrame( parsed_server_data )
```



# Web Server Logs

request_first_line	request_header_referer	request_header_user_agent	request_http_ver	request_method	request_url
GET /linux/doing-pxe-without-dhcp-control HTTP...	http://howto.basjes.nl/	Mozilla/5.0 (Windows NT 5.1; rv:35.0) Gecko/20...	1.1	GET	/linux/doing-pxe-v
GET /join_form HTTP/1.0	http://howto.basjes.nl/	Mozilla/5.0 (Windows NT 5.1; rv:35.0) Gecko/20...	1.0	GET	/join_form
POST /join_form HTTP/1.1	http://howto.basjes.nl/join_form	Mozilla/5.0 (Windows NT 5.1; rv:35.0) Gecko/20...	1.1	POST	/join_form
GET /join_form HTTP/1.0	http://howto.basjes.nl/	Mozilla/5.0 (Windows NT 6.3; WOW64; rv:34.0) G...	1.0	GET	/join_form
POST /join_form HTTP/1.1	http://howto.basjes.nl/join_form	Mozilla/5.0 (Windows NT 6.3; WOW64; rv:34.0) G...	1.1	POST	/join_form
GET /acl_users/credentials_cookie_auth/require...	http://howto.basjes.nl/join_form	Mozilla/5.0 (Windows NT 6.3; WOW64; rv:34.0) G...	1.1	GET	/acl_users/creden



# Windows Event Logs

```
pip install python-evtx
```

```
import Evtx.Evtx as evtx

xml = "<Events>"
with evtx.Evtx("window_event_log.evtx") as log:
    for record in log.records():
        xml += record.xml()
xml += "</Events>"
```





# Nested Data?





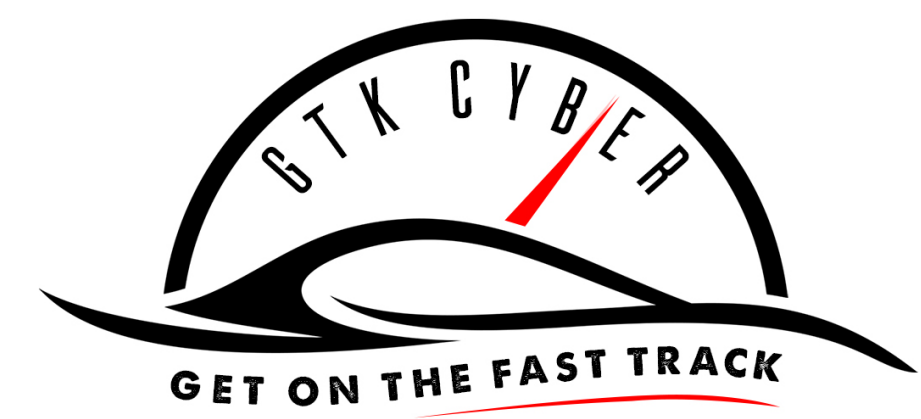
# Nested Data?

```
{ "time": 1084443427.311224,  
  "timestamp": "2004-05-13T10:17:07.311224",  
  "IP": {  
    "version": 4,  
    "ttl": 128,  
    "proto": 6,  
    "options": [],  
    "len": 48,  
    "dst": "65.208.228.223",  
    "frag": 0,  
    "flags": 2, "src": "145.254.160.237",  
    "chksum": 37355  
  },  
  "Ethernet": { "src": "00:00:01:00:00:00", "type": 2048, "dst": "fe:ff:20:00:01:00" },  
  ...  
}
```



# Nested Data

```
pd.read_json( 'nested_data.json' )
```



# Nested Data

```
pd.read_json( 'nested_data.json' )
```

	DNS	Ethernet	IP	TCP	UDP	time	timestamp
0	NaN	{'type': 2048, 'dst': 'fe:ff:20:00:01:00', 'sr...	{'dst': '65.208.228.223', 'len': 48, 'version'...	{'window': 8760, 'chksum': 49932, 'sport': 337...	NaN	1.084443e+09	2004-05-13 10:17:07.311224
1	NaN	{'type': 2048, 'dst': '00:00:01:00:00:00', 'sr...	{'dst': '145.254.160.237', 'len': 48, 'version'...	{'window': 5840, 'chksum': 23516, 'sport': 80,...	NaN	1.084443e+09	2004-05-13 10:17:08.222534
2	NaN	{'type': 2048, 'dst': 'fe:ff:20:00:01:00', 'sr...	{'dst': '65.208.228.223', 'len': 40, 'version'...	{'window': 9660, 'chksum': 31076, 'sport': 337...	NaN	1.084443e+09	2004-05-13 10:17:08.222534
3	NaN	{'type': 2048, 'dst': 'fe:ff:20:00:01:00', 'sr...	{'dst': '65.208.228.223', 'len': 519, 'version'...	{'window': 9660, 'chksum': 43352, 'sport': 337...	NaN	1.084443e+09	2004-05-13 10:17:08.222534
4	NaN	{'type': 2048, 'dst': '00:00:01:00:00:00', 'sr...	{'dst': '145.254.160.237', 'len': 40, 'version'...	{'window': 6432, 'chksum': 33825, 'sport': 80,...	NaN	1.084443e+09	2004-05-13 10:17:08.783340



# Nested Data

```
from pandas.io.json import json_normalize
```

```
import json
```

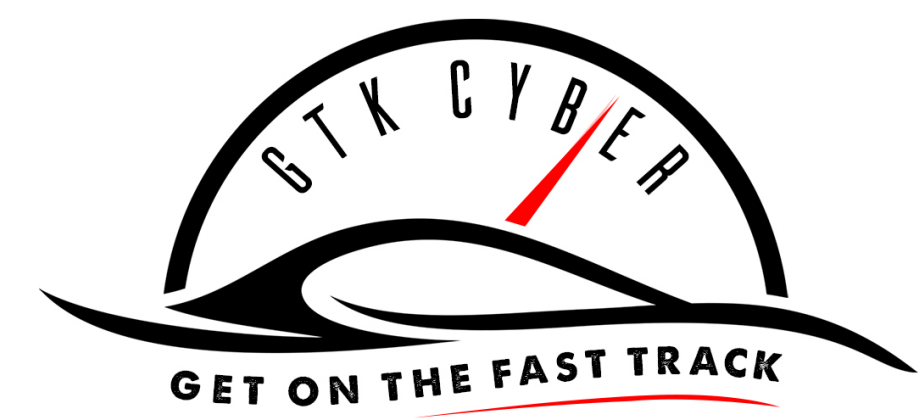
```
import pandas as pd
```

```
with open('nested.json') as data_file:
```

```
    pcap_data = json.load(data_file)
```

```
df = pd.DataFrame( json_normalize(pcap_data) )
```





# Nested Data

```
df = pd.DataFrame( json_normalize(pcap_data) )
```

...	TCP.seq	TCP.sport	TCP.urgptr	TCP.window	L
...	951057939.0	3372.0	0.0	8760.0	
...	290218379.0	80.0	0.0	5840.0	
...	951057940.0	3372.0	0.0	9660.0	
...	951057940.0	3372.0	0.0	9660.0	
...	290218380.0	80.0	0.0	6432.0	

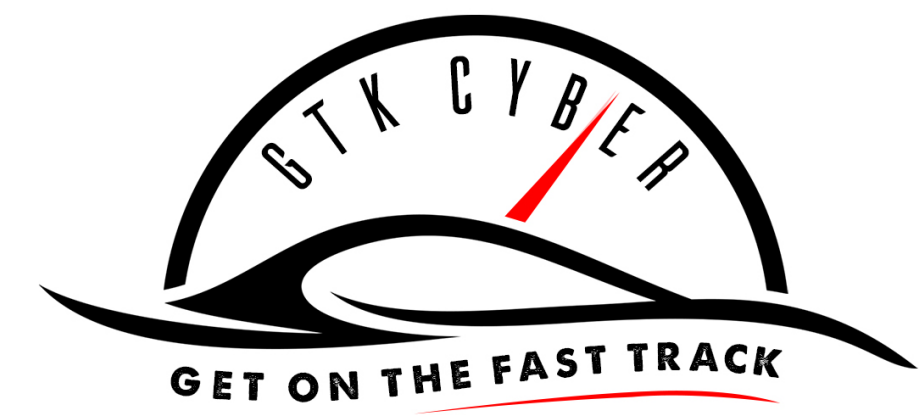


# ElasticSearch & Splunk



elasticsearch

splunk<sup>®</sup>>



# ElasticSearch & Splunk

```
pip install huntlib
```

```
e = ElasticDF(  
    url="https://localhost:9200",  
    ssl=True,  
    username="myuser",  
    password="mypass"  
)  
df = e.search_df(  
    lucene="proto:tcp AND port:80",  
    index="myindex-*",  
    days=1,  
    normalize=False  
)
```

```
s = SplunkDF(  
    host "<splunk_ip>",  
    username="myuser",  
    password="mypass"  
)  
df = s.search_df(  
    spl='search index=win_events src="10.9.*.*"',  
    start_time=datetime.now()-timedelta(days=2),  
    end_time=datetime.now()  
)
```



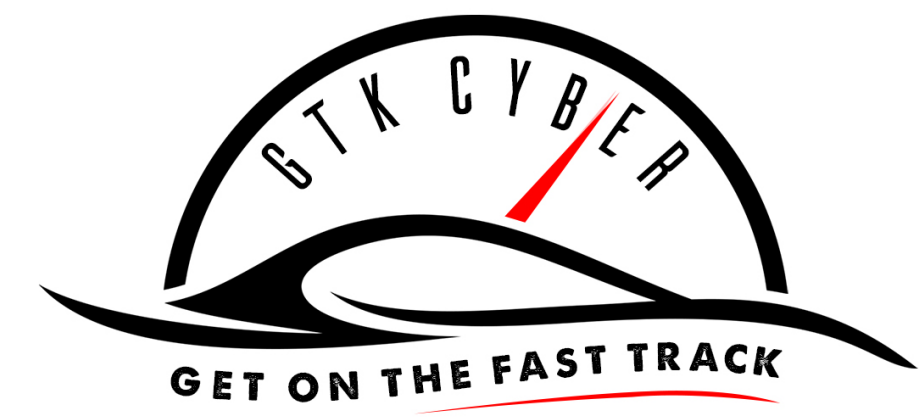
# Manipulating a DataFrame



```
series = df[ 'column1' ]
```

Returns a **series**

```
df[ 'ip' ] .value_counts() .head()
```



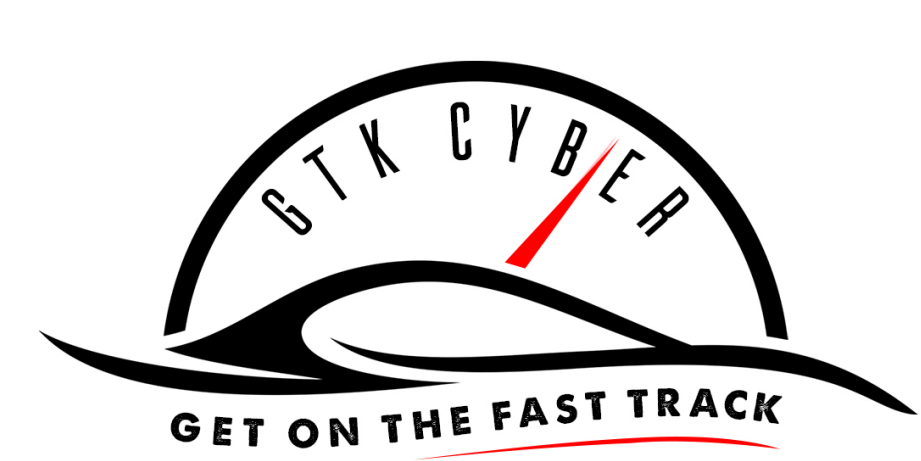
# Two ways of accessing columns

```
series = df['column1']
```

```
series = df.column1
```

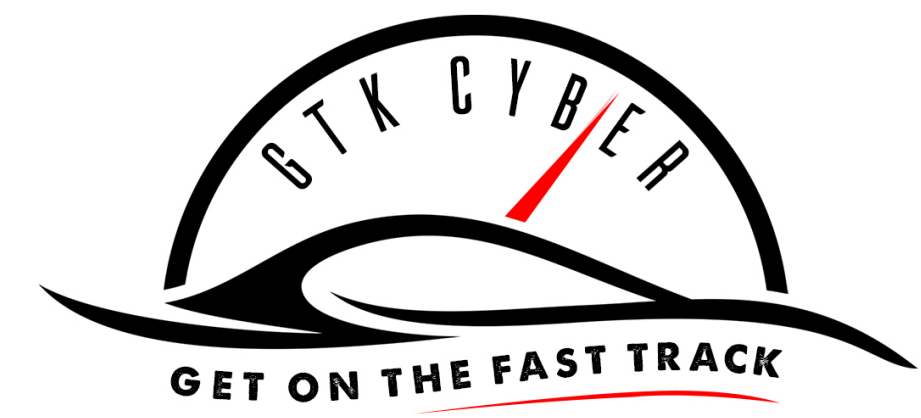
Don't use the dots!





```
df = df[ ['column1', 'column2', 'column3'] ]
```

Returns a DataFrame



# Filtering a DataFrame

```
df[<boolean condition>]
```

```
df[['col1', 'col2']][df['col3'] > 5]
```

Columns

Filter





```
data.loc[<index>]  
data.loc[<list of indexes>]  
data.sample(<n>)
```



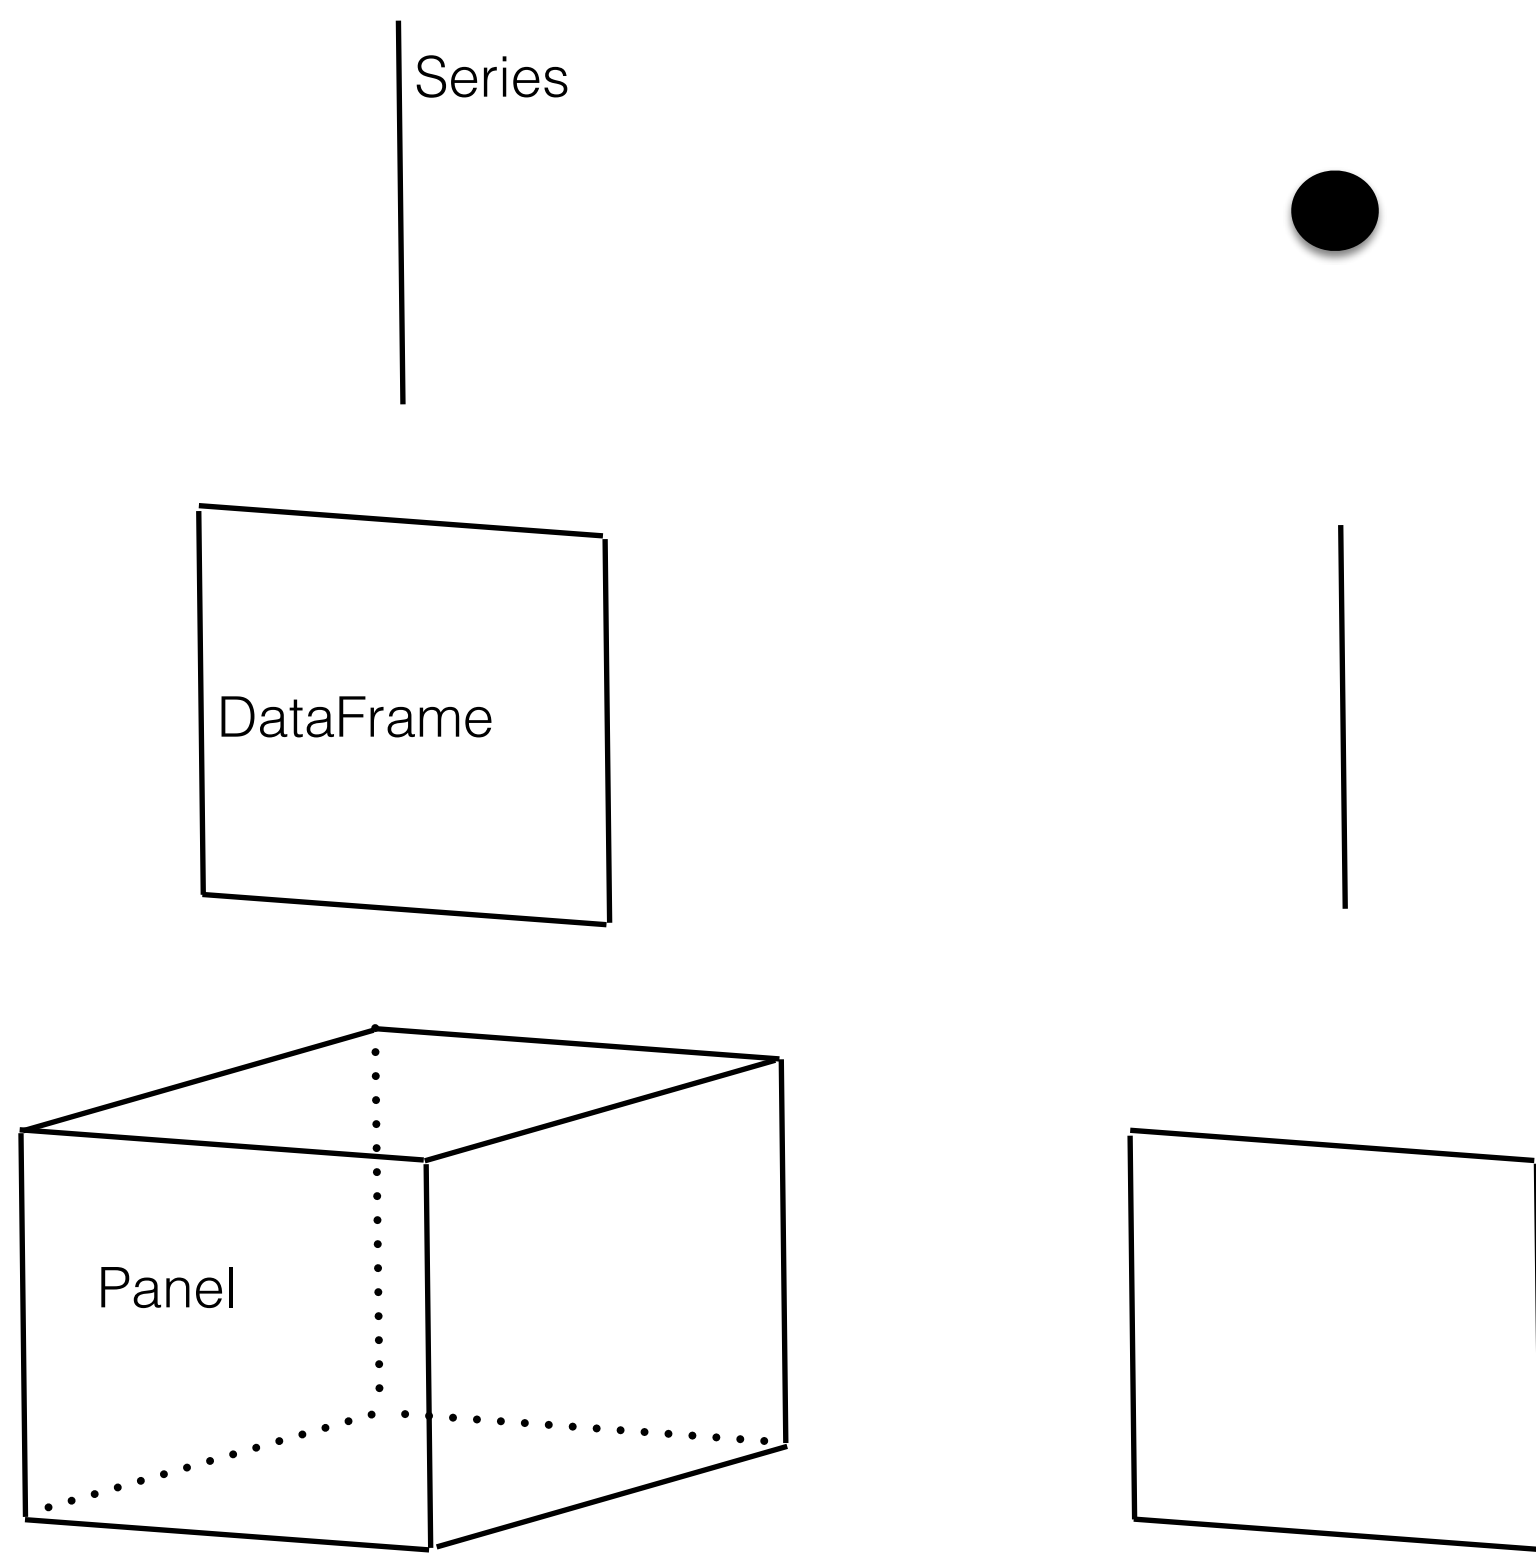
data.apply( <function> )

IF

Call Apply On...

THEN


Function  
Receives






data.apply( <function> )


**IF**  
Call Apply On...




**THEN**  
Function  
Receives



**IF**  
Function  
Returns



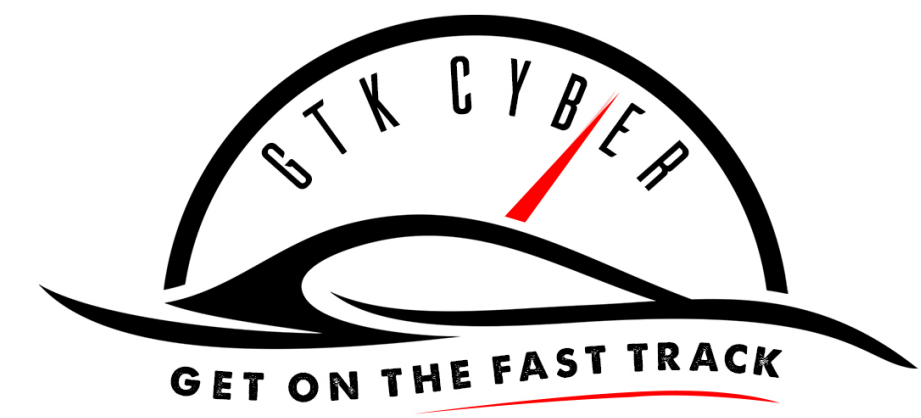
**THEN**  
Apply  
Returns...





# In Class Exercise:

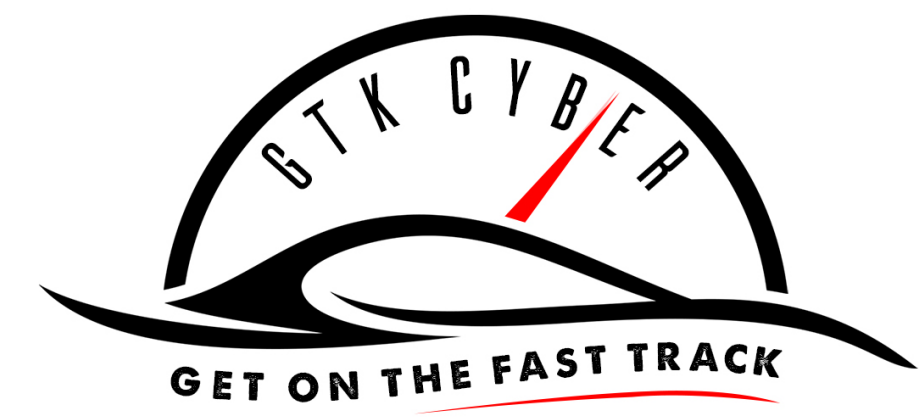
Please complete Worksheet 2.1: Exploring Two Dimensional Data



# Exercise 1

```
#Data1
{"first_name":{"0":"Robert","1":"Steve","2":"Anne","3":"Alice"},
"last_name":{"0":"Hernandez","1":"Smith","2":"Raps","3":"Muller"},
"birthday":{"0":"5\\/3\\/67","1":"8\\/4\\/84","2":"9\\/13\\/
91","3":"4\\/15\\/75"}
}
```

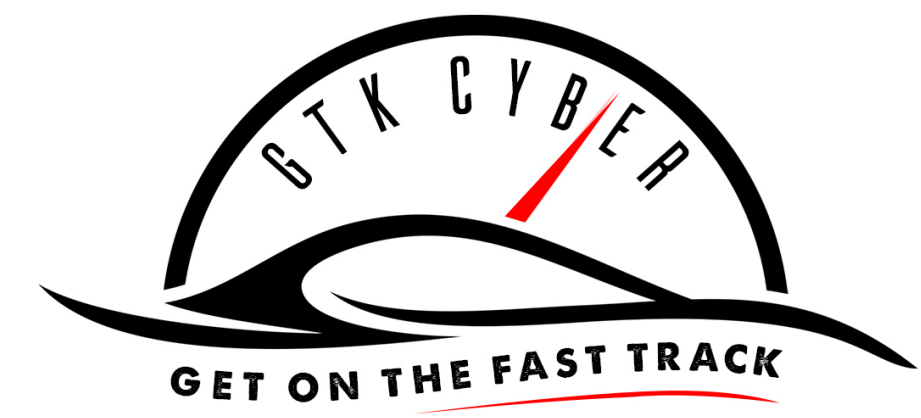
```
df1 = pd.read_json( '../data/data1.json' )
```



# Exercise 1

```
#Data2
{"0":{"first_name":"Robert","last_name":"Hernandez","birthday":"5\\/3\\
\\/67"},
"1":{"first_name":"Steve","last_name":"Smith","birthday":"8\\/4\\
84"},
"2":{"first_name":"Anne","last_name":"Raps","birthday":"9\\/13\\/91"},
"3":{"first_name":"Alice","last_name":"Muller","birthday":"4\\/15\\
75"}}
```


```
df2 = pd.read_json( '.. /data/data2.json', orient='index' )
```



# Exercise 1

```
#Data3
[
{"first_name": "Robert", "last_name": "Hernandez", "birthday": "5\\/3\\/67"},
{"first_name": "Steve", "last_name": "Smith", "birthday": "8\\/4\\/84"},
{"first_name": "Anne", "last_name": "Raps", "birthday": "9\\/13\\/91"},
{"first_name": "Alice", "last_name": "Muller", "birthday": "4\\/15\\/75"}]
```

```
df3 = pd.read_json( '../data/data3.json',
orient='columns' )
```



# Exercise 1

#Data4

```
{ "columns": [ "first_name", "last_name", "birthday" ],  
  "index": [ 0, 1, 2, 3 ],  
  "data": [ [ "Robert", "Hernandez", "5\\/3\\/67" ], [ "Steve", "Smith", "8\\/4\\/84" ], [ "Anne", "Raps", "9\\/13\\/91" ], [ "Alice", "Muller", "4\\/15\\/75" ] ] }
```

```
df4 = pd.read_json( '.. /data/data4.json', orient='split' )
```





# Exercise 2

#Write the functions

```
def get_os(x):  
    user_agent = parse(x)  
    return user_agent.os.family
```

```
def get_browser(x):  
    user_agent = parse(x)  
    return user_agent.browser.family
```



# Exercise 2

```
#Apply the functions to the dataframe
server_df['os'] =
server_df['request_header_user_agent'].apply( get_os )
server_df['browser'] =
server_df['request_header_user_agent'].apply( get_browser )
```



# Exercise 2

```
#Apply the functions to the dataframe
server_df['os'] =
server_df['request_header_user_agent'].apply( get_os )
server_df['browser'] =
server_df['request_header_user_agent'].apply( get_browser )
```

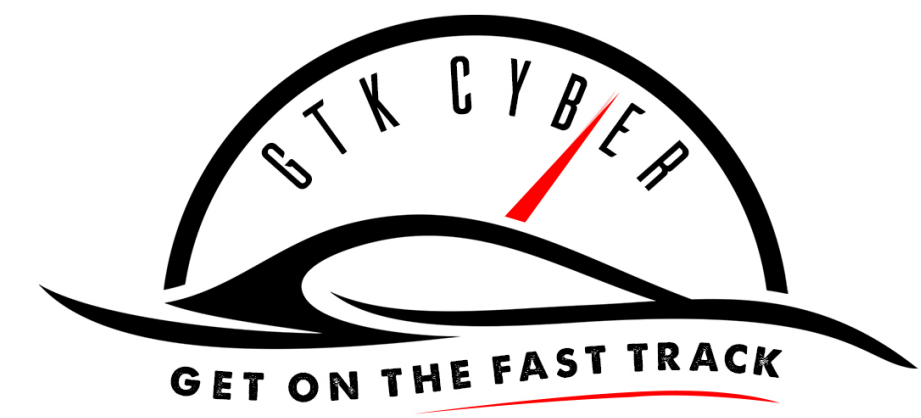
```
#Get the top 10 values
server_df['os'].value_counts().head(10)
```



# Exercise 2

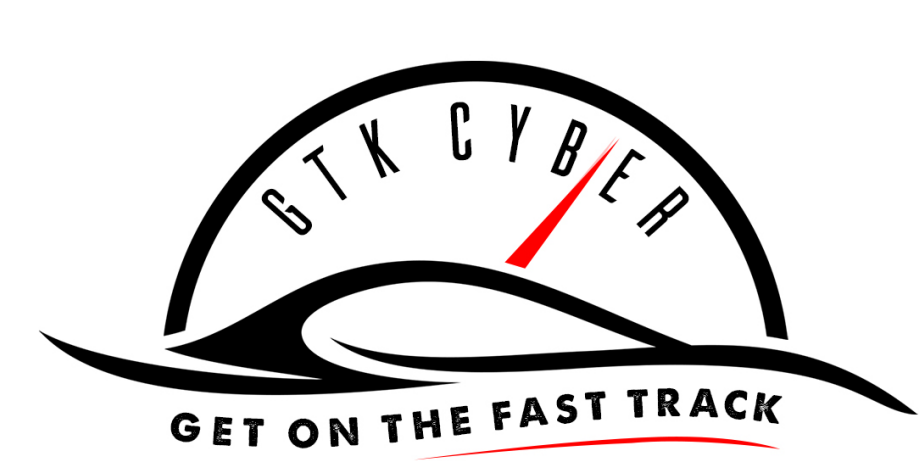
```
#Get the top 10 values  
server_df['os'].value_counts().head(10)
```

Windows 7	2041
Windows Vista	500
Windows XP	423
Windows 8.1	221
Linux	125
Mac OS X	80
Chrome OS	60
Ubuntu	6

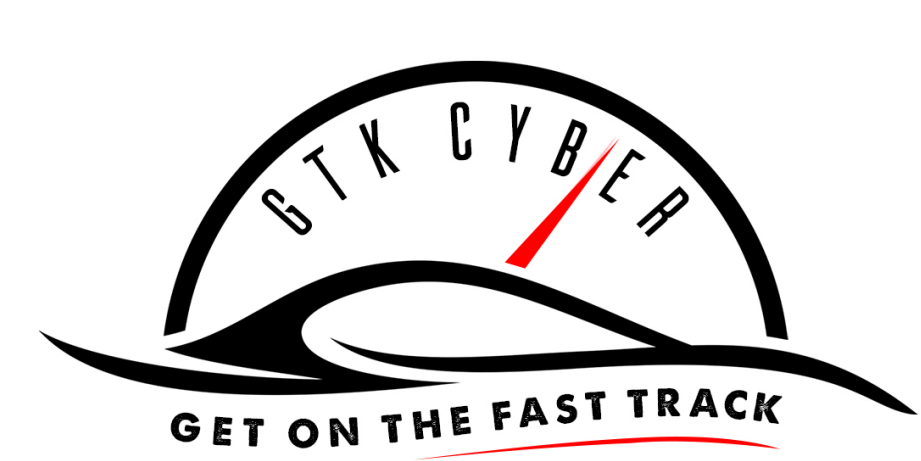


# Exercise 3

```
bots = pd.read_csv( '../data/dailybots.csv' )  
gov_bots = bots[['botfam', 'hosts']][bots['industry'] ==  
"Government/Politics"]  
  
gov_bots.groupby( 'botfam', as_index=False ).sum( )
```



# Questions?

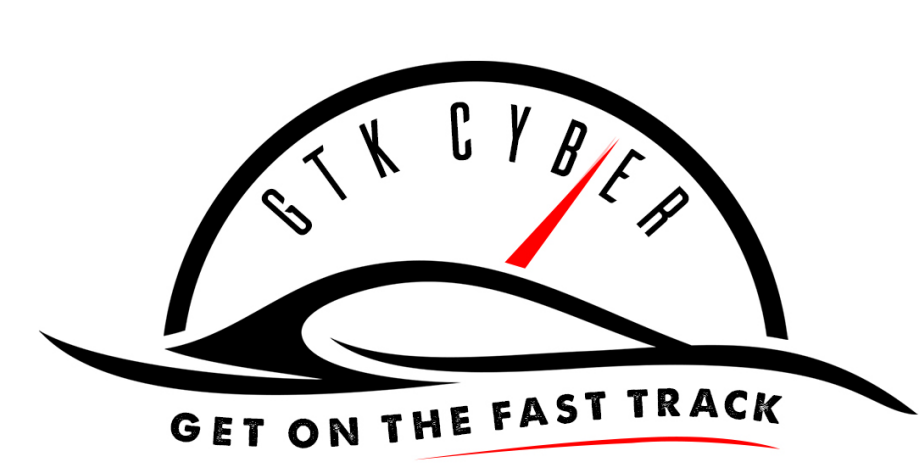


1	2	3
4	5	6
7	8	9

data.T

1	4	7
2	5	8
3	6	9





*#Gets you the sum of columns*  
`data.sum( axis=0 )`

*#Gets you the sum of the rows*  
`data.sum( axis=1 )`



```
DataFrame.drop(labels,  
axis=0,  
level=None,  
inplace=False,  
errors='raise')¶
```



# Merging Data Sets



# Union



# Union

Series 1

Index	Value
0	6
1	4
2	2
3	3





# Union

Series 1

Index	Value
0	6
1	4
2	2
3	3

Series 2

Index	Value
0	7
1	5
2	3
3	4



# Union

combinedSeries

Index	Value
0	6
1	4
2	2
3	3
4	7
5	5
6	3
7	4




```
combinedSeries = pd.concat(  
    [series1, series2],  
    ignore_index=True )
```



# Joins



# Inner Join (Intersection)




Inner Join






# Left Join



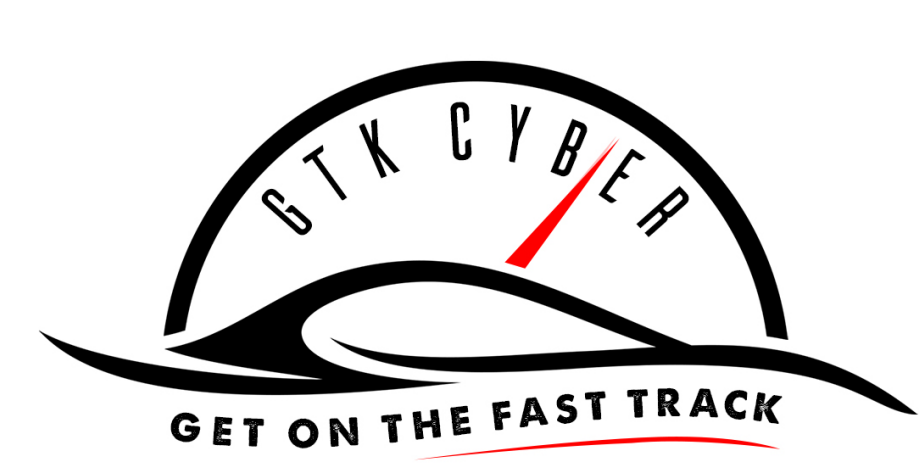
Left Join



# Right Join




Right Join



```
pd.merge( leftData, rightData )
```

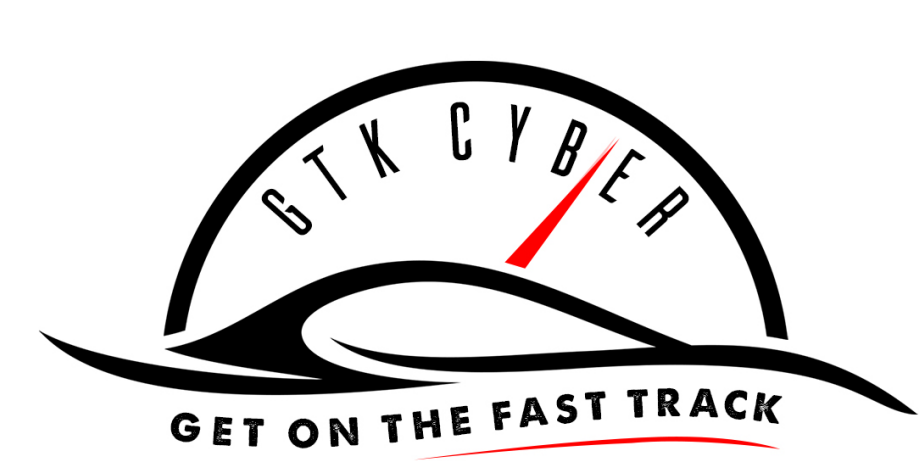


```
pd.merge( leftData,  
          rightData,  
          how="<join type>" )
```



```
pd.merge( leftData, rightData,  
how="<join type>" )
```

Option	SQL Equivalent	Description
inner	INNER JOIN	Intersection
left	LEFT OUTER JOIN	Returns items in Set A, but not in Set B
right	RIGHT OUTER JOIN	Returns items in Set B, but not in Set A
outer	FULL OUTER JOIN	Returns the union of both sets



```
pd.merge( leftData, rightData,  
          how="<join type>",  
          on="<field list> )
```





# Grouping and Aggregating Data



# Grouping and Aggregating Data

date	src_ip	dst_ip	port
2018-06-21	192.168.20.2	10.10.4.1	80
2018-06-21	192.168.20.1	10.10.4.2	443
2018-06-21	192.168.20.2	10.10.5.1	80
2018-06-22	192.168.20.2	10.10.4.1	80

**`df.groupby(field or list of fields)`**

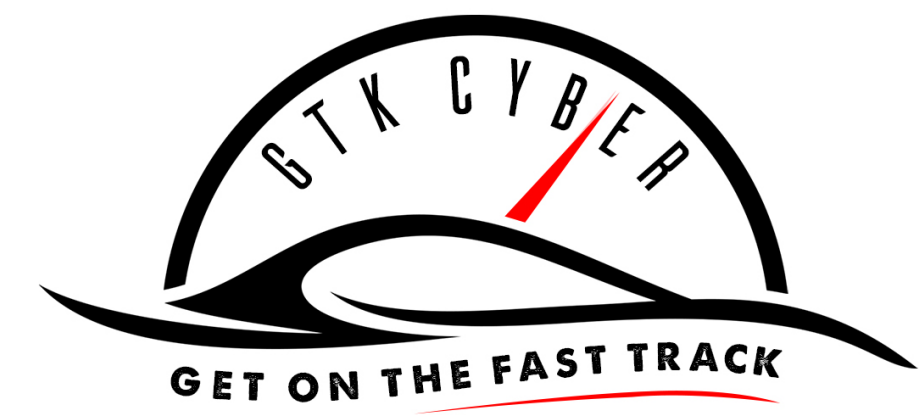


# Grouping and Aggregating Data

date	src_ip	dst_ip	port
2018-06-21	192.168.20.2	10.10.4.1	80
2018-06-21	192.168.20.1	10.10.4.2	443
2018-06-21	192.168.20.2	10.10.5.1	80
2018-06-22	192.168.20.2	10.10.4.1	80

```
df.groupby( 'src_ip' )[ 'port' ].count( )
```

src_ip	count
192.168.20.1	1
192.168.20.2	3




# Grouping and Aggregating Data

date	src_ip	dst_ip	port
2018-06-21	192.168.20.2	10.10.4.1	80
2018-06-21	192.168.20.1	10.10.4.2	443
2018-06-21	192.168.20.2	10.10.5.1	80
2018-06-22	192.168.20.2	10.10.4.1	80

```
df.groupby(['date', 'src_ip'])['port'].count()
```

Multi-Index

date	src_ip	
2018-06-21	192.168.20.1	1
	192.168.20.2	2
2018-06-22	192.168.20.2	1



# Grouping and Aggregating Data

date	src_ip	dst_ip	port
2018-06-21	192.168.20.2	10.10.4.1	80
2018-06-21	192.168.20.1	10.10.4.2	443
2018-06-21	192.168.20.2	10.10.5.1	80
2018-06-22	192.168.20.2	10.10.4.1	80

```
df.groupby(['date', 'src_ip'], as_index=False)  
['port'].count()
```

	date	src_ip	port
0	2018-06-21	192.168.20.1	1
1	2018-06-21	192.168.20.2	2
2	2018-06-22	192.168.20.2	1

# Grouping Functions

## Pivot

df

	foo	bar	baz	zoo
0	one	A	1	x
1	one	B	2	y
2	one	C	3	z
3	two	A	4	q
4	two	B	5	w
5	two	C	6	t



```
df.pivot(index='foo',
          columns='bar',
          values='baz')
```

bar	A	B	C
foo			
one	1	2	3
two	4	5	6



# Grouping Functions

## Melt

df3

	first	last	height	weight
0	John	Doe	5.5	130
1	Mary	Bo	6.0	150




df3.melt(id\_vars=['first', 'last'])

	first	last	variable	value
0	John	Doe	height	5.5
1	Mary	Bo	height	6.0
2	John	Doe	weight	130
3	Mary	Bo	weight	150


# Grouping Functions

## Stack



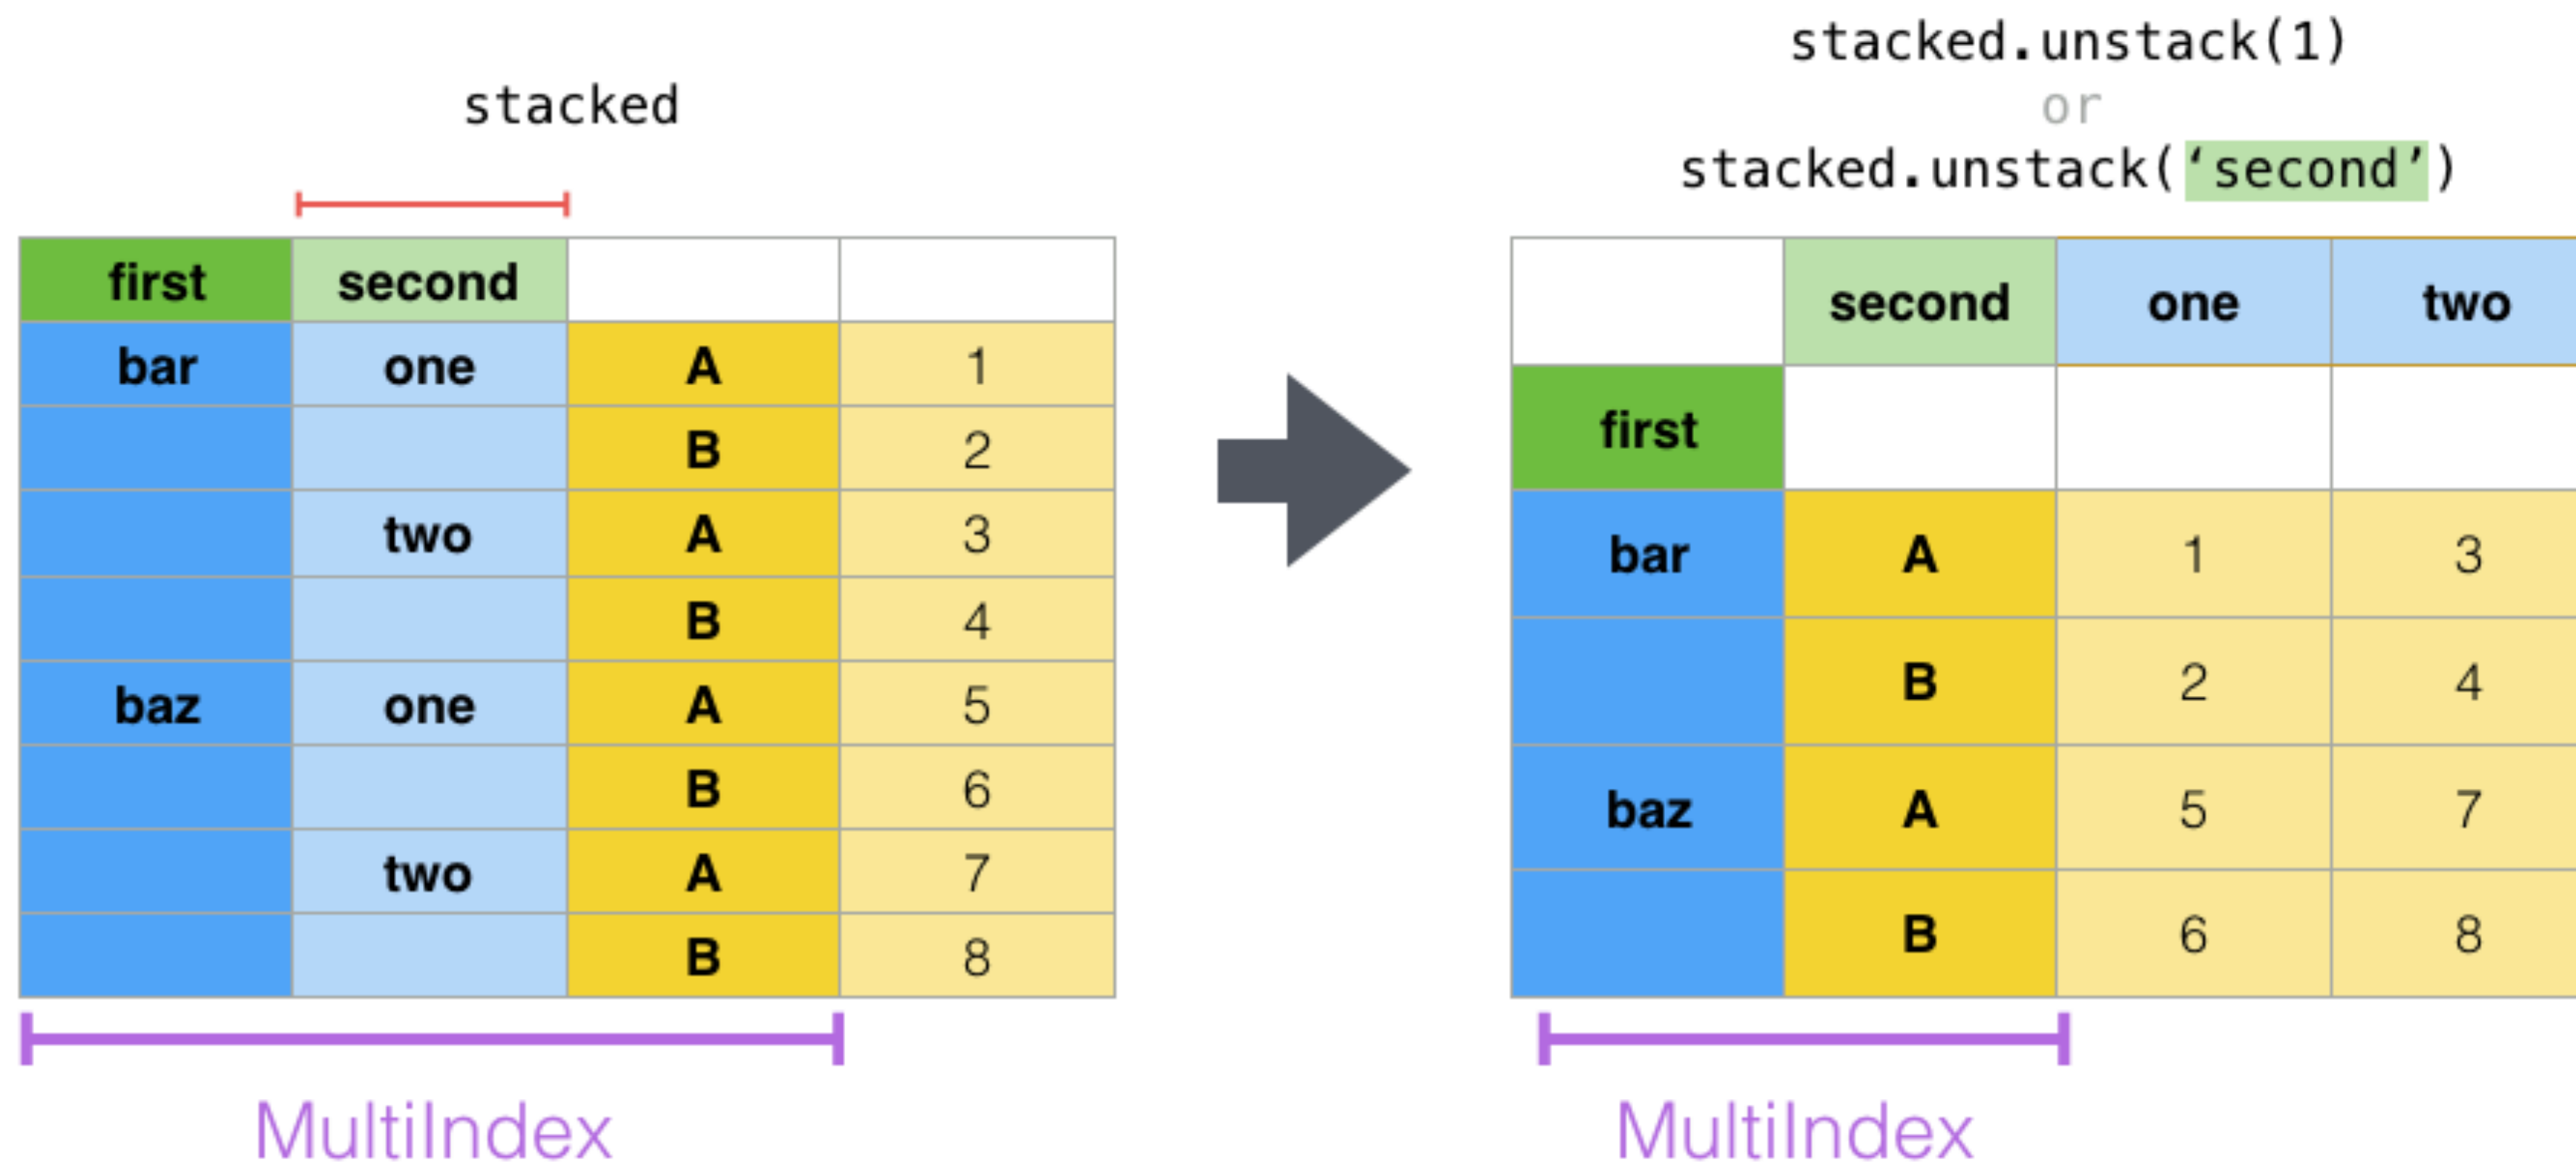
# Grouping Functions

## Unstack



# Grouping Functions

## Unstack(1)





```
Series.dropna( )
```

```
Series.fillna(value="<something>" )
```





# Another Approach: Drill



# Another Approach: Drill

- Drill is an open source query engine which allows you to query many kinds of self-describing data using ANSI SQL.
- Drill is fast, scalable and extremely versatile.







# Another Approach: Drill

Out of the box, Drill can query:

- Delimited Data (csv, tsv, psv)
- Apache log files
- Avro
- Parquet
- JSON
- PCAP
- Syslog
- And more...

There are extensions available on GitHub that allow you to query log files and other data.



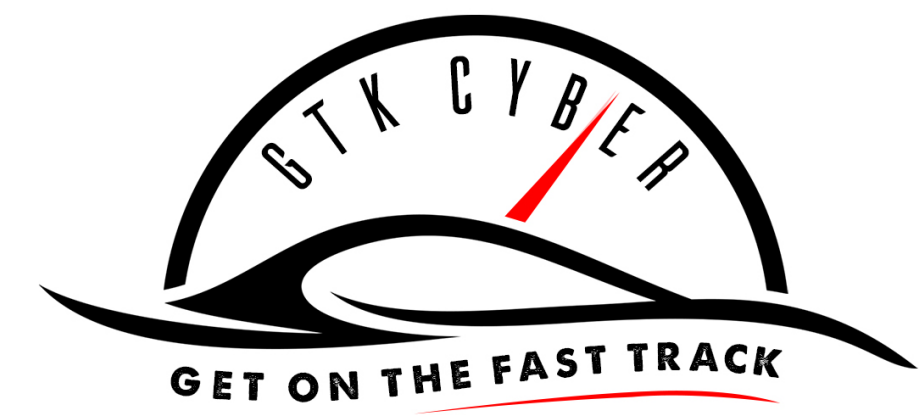


# Another Approach: Drill

Out of the box, Drill can connect to:

- MySQL (or any JDBC compliant database)
- Hadoop
- Kafka
- S3/Azure/Google Cloud
- HBase
- Hive
- MongoDB
- And others...





# Another Approach: Drill

```
from pydrill.client import PyDrill

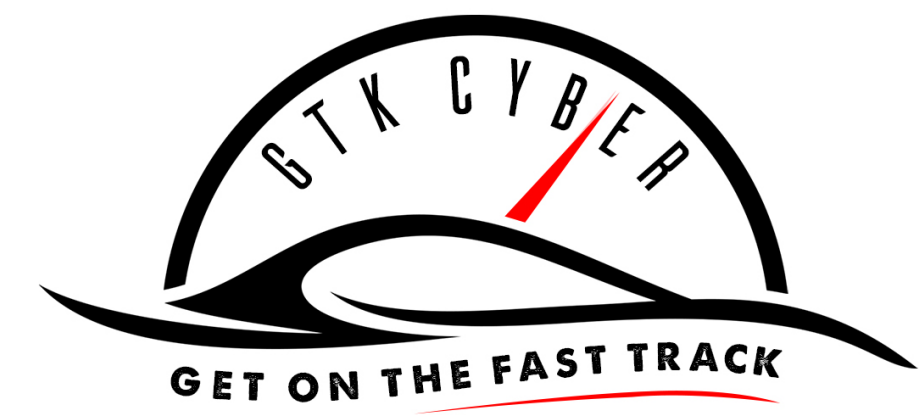
drill = PyDrill(host='localhost', port=8047)

if not drill.is_active():
    raise ImproperlyConfigured('Please run Drill first')

query = drill.query('' '<Your query here>'')

df = query.to_dataframe()
```





# Another Approach: Drill

```
SELECT src_port, dst_port, COUNT(*) AS packet_count  
FROM dfs.test.`dns.pcap`  
GROUP BY src_port, dst_port  
ORDER BY packet_count DESC
```





# In-Class Exercise

Please complete Worksheet 2.2: Exploratory Data Analysis